




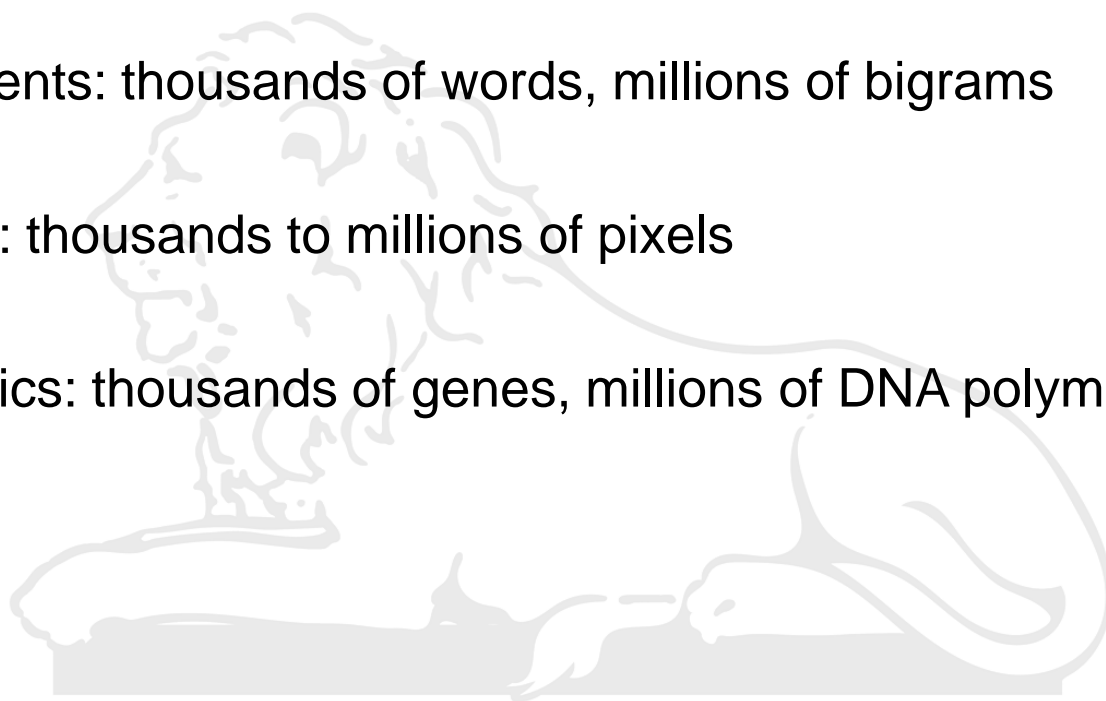
PCA & LDA

Dimensionality Reduction in ML



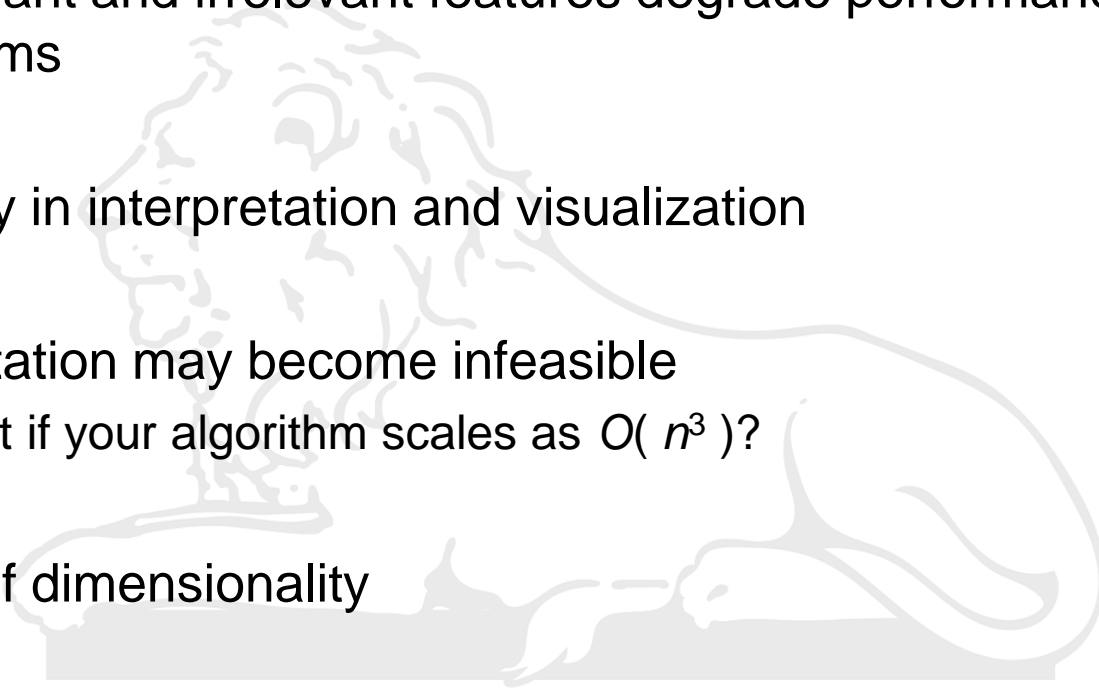
Dimensionality Reduction

- Many modern data domains involve huge numbers of features / dimensions
 - Documents: thousands of words, millions of bigrams
 - Images: thousands to millions of pixels
 - Genomics: thousands of genes, millions of DNA polymorphisms



Why to reduce dimensions?

- High dimensionality has many costs
 - Redundant and irrelevant features degrade performance of some ML algorithms
 - Difficulty in interpretation and visualization
 - Computation may become infeasible
 - what if your algorithm scales as $O(n^3)$?
 - Curse of dimensionality



A Toy Example

	Educations	Jobs opportunities	Safety	Air quality	Entertainment	Connectivity	Foods
City-1	8	7	7	9	8	7	8
City-2	6	8	7	6	8	9	8
City-3	9	10	7	8	8	4	8
City-4	4	4	6	2	8	5	7
City-5	3	6	6	5	8	9	8
City-6	2	7	6	3	8	4	7
City-7	7	8	7	8	8	6	8
City-8	8	3	6	3	8	8	7
City-9	5	9	7	7	8	5	7
City-10	6	8	6	9	8	9	8

Another example

	Educations	Jobs opportunities	Safety	Air quality	Entertainment	Connectivity	Foods
City-1	8	7	7	9	8	7	8
City-2	6	8	7	6	8	9	8
City-3	9	10	7	8	8	4	8
City-4	4	4	6	2	8	5	7
City-5	3	6	6	5	8	9	8
City-6	2	7	6	3	8	4	7
City-7	7	8	7	8	8	6	8
City-8	8	3	6	3	8	8	7
City-9	5	9	7	7	8	5	7
City-10	6	8	6	9	8	9	8

Dimension reduction

■ Two approaches to perform dim. reduction $\mathbb{R}^N \rightarrow \mathbb{R}^M$ ($M < N$)

- **Feature selection:** choosing a subset of all the features

$$[x_1 \ x_2 \dots x_N] \xrightarrow{\text{feature selection}} [x_{i_1} \ x_{i_2} \dots x_{i_M}]$$

- **Feature extraction:** creating new features by combining existing ones

$$[x_1 \ x_2 \dots x_N] \xrightarrow{\text{feature extraction}} [y_1 \ y_2 \dots y_M] = f([x_{i_1} \ x_{i_2} \dots x_{i_M}])$$

- In either case, the goal is to find a low-dimensional representation of the data that preserves (most of) the information or structure in the data

■ Linear feature extraction

- The “optimal” mapping $y=f(x)$ is, in general, a non-linear function whose form is problem-dependent
 - Hence, feature extraction is commonly limited to linear projections $\mathbf{y}=\mathbf{W}\mathbf{x}$

$$\begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_N \end{bmatrix} \xrightarrow{\text{linear feature extraction}} \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_M \end{bmatrix} = \begin{bmatrix} w_{11} & w_{12} & \dots & w_{1N} \\ w_{21} & w_{22} & \dots & w_{2N} \\ \vdots & \vdots & \ddots & \vdots \\ w_{M1} & w_{M2} & \dots & w_{MN} \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_N \end{bmatrix}$$

- Linear methods
 - Principal component analysis (PCA)
 - Multidimensional scaling (MDS)
 - Independent component analysis (ICA)
- Nonlinear methods
 - Kernel PCA
 - Locally linear embedding (LLE)
 - Laplacian eigenmaps (LEM)
 - Semidefinite embedding (SDE)

Statistics Preliminaries

■ Mean

- Let X_1, X_2, \dots, X_n be n observations of a random variable X

$$\mu = \bar{X} = E[X] = \frac{1}{n} \sum_{i=1}^n X_i$$

- Mean is a measure of central tendency (others are mode and median)

■ Standard Deviation

- Measure of variability (square root of variance)

$$\sigma = \sqrt{\frac{1}{n} \sum_{i=1}^n (X_i - \mu)^2}$$

Statistics Preliminaries

■ Covariance

- A measure of how two variables change together

$$\Sigma_{XY} = \frac{1}{n} \sum_{i=1}^n (X_i - \mu_X)(Y_i - \mu_Y)^T$$


$$\Sigma_{XY} = E[(X - E[X])(Y - E[Y])]$$

■ Covariance Matrix

$$\Sigma = \begin{pmatrix} \sigma_{X1}^2 & \Sigma_{X1X2} & \dots & \Sigma_{X1Xd} \\ \Sigma_{X1X2} & \sigma_{X2}^2 & \dots & \Sigma_{X2Xd} \\ \cdot & & \dots & \cdot \\ \cdot & & & \cdot \\ \cdot & \Sigma_{X1Xd} & \Sigma_{X2Xd} & \dots & \sigma_{Xd}^2 \end{pmatrix}$$

Statistics Preliminaries – An Example


$$\mu = \begin{pmatrix} 2 \\ 3 \end{pmatrix} \quad \Sigma = \begin{pmatrix} 2 & 0 \\ 0 & 3 \end{pmatrix}$$



PCA – An Example

$$\mu = \begin{pmatrix} 2 \\ 3 \end{pmatrix} \quad \Sigma = \begin{pmatrix} 2 & 2 \\ 2 & 3 \end{pmatrix}$$


Can you find a vector that would approximate this 2-D space?



PCA – Let's build some intuition

$$\mu = \begin{pmatrix} 2 \\ 3 \end{pmatrix} \quad \Sigma = \begin{pmatrix} 2 & 2 \\ 2 & 3 \end{pmatrix}$$


*Can you find a vector that
would approximate this 2-D space?*
Green vector right?



PCA – Let's build some intuition

$$\mu = \begin{pmatrix} 0 \\ 0 \end{pmatrix} \quad \Sigma = \begin{pmatrix} 2 & 2 \\ 2 & 3 \end{pmatrix}$$

*It would be nice to
diagonalize the covariance matrix
then you have only think about variance*
Think eigenvectors of covariance matrix




A 2-D Example

Step-1

<http://kybele.psych.cornell.edu/~edelman/Psych-465-Spring-2003/PCA-tutorial.pdf>

- DATA:

x	y
2.5	2.4
0.5	0.7
2.2	2.9
1.9	2.2
3.1	3.0
2.3	2.7
2	1.6
1	1.1
1.5	1.6
1.1	0.9



mean

this becomes the
new origin of the
data from now on

A 2-D Example

Step-2

- Calculate the covariance matrix

$$\text{cov} = \begin{pmatrix} .616555556 & .615444444 \\ .615444444 & .716555556 \end{pmatrix}$$

Step-3

Calculate the eigenvectors and eigenvalues of the covariance matrix

$$\text{eigenvalues} = \begin{pmatrix} .0490833989 \\ 1.28402771 \end{pmatrix}$$

$$\text{eigenvectors} = \begin{pmatrix} -.735178656 & -.677873399 \\ .677873399 & -.735178656 \end{pmatrix}$$

A 2-D Example

Step-4

- Feature Vector

FeatureVector = (eig₁ eig₂ eig₃ ... eig_n)

We can either form a feature vector with both of the eigenvectors:

$$\begin{pmatrix} -.677873399 & -.735178656 \\ -.735178656 & .677873399 \end{pmatrix}$$

or, we can choose to leave out the smaller, less significant component and only have a single column:


$$\begin{pmatrix} -.677873399 \\ -.735178656 \end{pmatrix}$$

Transformed Data (Single eigenvector)

x


-.827970186
 1.77758033
 -.992197494
 -.274210416
 -1.67580142
 -.912949103
 .0991094375
 1.14457216
 .438046137
 1.22382056

Original PCA data



2D point cloud


Original data restored using only a single eigenvector



Approximation using
one eigenvector basis

Limitations of PCA?

- PCA projects the data in lower dimensional space spanned by the directions of maximum variance (most accurate *data*)




- However the directions of maximum variance may be useless for classification (linear...). LDA overcomes this limitation of PCA.


LDA → Lower dimensions + preserve classification property

Bias

- **LDA projects the data in lower dimensional space in such a way that patterns from different classes are well separated.**
- How to find such a line (in 2D), and lower dimensional space in general?



Projection of a feature vector on to a space



LDA:

Input: We have a binary classification problem, where patterns are coming from a d -dimensional space.

→ n_1 samples are from class-1.

→ n_2 samples are from class-2.

$$n_1 + n_2 = n$$

Samples are given by vectors

$$x_1, x_2, x_3, \dots, x_n$$

LDA...

- * Let μ_1 and μ_2 be the means of patterns from class-1 and class-2, respectively in original space.
- * Moreover, $\tilde{\mu}_1$ and $\tilde{\mu}_2$ are mean of patterns after projection.

Then,

$$\tilde{\mu}_1 = \frac{1}{n_1} \sum_{x_i \in C_1} v^T x_i = v^T \left(\frac{1}{n_1} \sum x_i \right) = \underline{v^T \mu_1}$$


$$\text{IIly } \underline{\tilde{\mu}_2 = v^T \mu_2}$$

LDA..

Our aim is to have the patterns of these two classes away from each other after projection.

So maximizing distance between their means may be a good measure:

→ The larger $|\tilde{\mu}_1 - \tilde{\mu}_2|$, the better is separation.




* The vertical axis is a better direction

* However,
 $|\hat{\mu}_1 - \hat{\mu}_2| > |\tilde{\mu}_1 - \tilde{\mu}_2|$

LDA

Therefore, we are missing something. Yes, we are not considering variance of the classes.



LDA...

Define the scatters as follows:

$$\tilde{S}_1^2 = \sum_{y_i \in C_1} (y_i - \tilde{M}_1)^2 \text{ for } \underline{\text{class-1}}$$

$$\text{and } \tilde{S}_2^2 = \sum_{y_i \in C_2} (y_i - \tilde{M}_2)^2 \text{ for } \underline{\text{class-2}}$$

$$\text{where, } y_i = V^T x_i$$

Then, we need to Maximize $\underline{|\tilde{M}_1 - \tilde{M}_2|}$
and minimize $\underline{\tilde{S}_1^2}$ and $\underline{\tilde{S}_2^2}$

Objective function

want projected means are far from each other

$$J(\mathbf{v}) = \frac{\overbrace{(\tilde{\mu}_1 - \tilde{\mu}_2)^2}}{\tilde{\mathbf{s}}_1^2 + \tilde{\mathbf{s}}_2^2}$$

want scatter in class 1 is as small as possible, i.e. samples of class 1 cluster around the projected mean $\tilde{\mu}_1$

want scatter in class 2 is as small as possible, i.e. samples of class 2 cluster around the projected mean $\tilde{\mu}_2$

LDA

Therefore,

Our objective function can be rewritten as

$$\arg \max_V J(V) = \frac{(\tilde{\mu}_1 - \tilde{\mu}_2)^2}{\tilde{S}_1^2 + \tilde{S}_2^2} = \frac{V^T S_B V}{V^T S_W V}$$

Now

$$\frac{dJ(V)}{dV} = 0 \Rightarrow S_B V - \frac{V^T S_B V (S_W V)}{V^T S_W V} = 0$$

$$\Rightarrow S_B V = \lambda S_W V$$

$$\Rightarrow \boxed{(S_W^{-1} S_B) V = \lambda V} \quad \text{Eigenvalue Problem}$$

Simplification

$$S_W^{-1} S_B V = \lambda V$$

* But $S_B X$ for any vector X , and $(\mu_1 - \mu_2)$ are parallel, i.e.

$$S_B X = (\mu_1 - \mu_2) \underbrace{(\mu_1 - \mu_2)^T X}_{\rightarrow \alpha} = \underline{\alpha (\mu_1 - \mu_2)}$$

Hence

$$V = S_W^{-1} (\mu_1 - \mu_2) \quad \text{---} \quad (*)$$

Summary:- From x_1, x_2, \dots, x_n , calculate S_W and $(\mu_1 - \mu_2) \Rightarrow$ Then $(*)$ gives you direction of the line.

Example

Consider the following data

$$C_1: [(1,2), (2,3), (3,3), (4,5), (5,5)]^T$$

$$C_2: [(1,0), (2,1), (3,1), (3,2), (5,3), (6,5)]^T$$

$$\text{Then } \mu_1 = [3 \quad 3.6]^T \text{ and } \mu_2 = [3.3 \quad 2]^T$$

Also,

$$S_1 = 4 * \text{Cov}(C_1) = \begin{bmatrix} 10 & 8 \\ 8 & 7.2 \end{bmatrix}$$

$$S_2 = 4 * \text{Cov}(C_2) = \begin{bmatrix} 17.3 & 16 \\ 16 & 16 \end{bmatrix}$$

$$S_W = S_1 + S_2 = \begin{bmatrix} 27.3 & 24 \\ 24 & 23.2 \end{bmatrix}$$


Example...

Hence, $S_{\bar{w}}^{-1} = \begin{bmatrix} .39 & -.41 \\ -.41 & .47 \end{bmatrix}$

Therefore, the optimal line direction V

$$V = S_{\bar{w}}^{-1} (M_1 - M_2) = \begin{bmatrix} -0.79 \\ 0.89 \end{bmatrix}$$

$$Y_1 = V^T C_1 \quad \text{and} \quad Y_2 = V^T C_2$$



THANKYOU

